



Databasteknik MongoDB

Utbildare: Mikael Olsson

mikael.olsson@emmio.se

076-174 90 43

NACKADEMIN

Socrative

<https://www.socrative.com/>

- Frågehanterare
 - Logga in som student
 - Ange rum “Emmio”
 - Få upp en vänta-skärm



Waiting for the next activity to begin...

Snabbfråga

- MySQL är en relationsdatabas
- MongoDB är en dokumentdatabas

Modelling

Users

ID	first_name	surname	cell	city	location_x	location_y
1	Paul	Miller	44755750561 1	London	45.123	47.232

Professions

ID	user_id	profession
10	1	banking
11	1	finance
12	1	trader

Cars

ID	user_id	model	year
20	1	Bentley	1973
21	1	Rolls Royce	1965

Modelling

- Det viktiga när man designar modeller i MongoDB är hur man ska använda datat, inte dess inbördes relationer.

Modelling

Relational

Users

ID	first_name	surname	cell	city	location_x	location_y
1	Paul	Miller	447557505611	London	45.123	47.232

Professions

ID	user_id	profession
10	1	banking
11	1	finance
12	1	trader

Cars

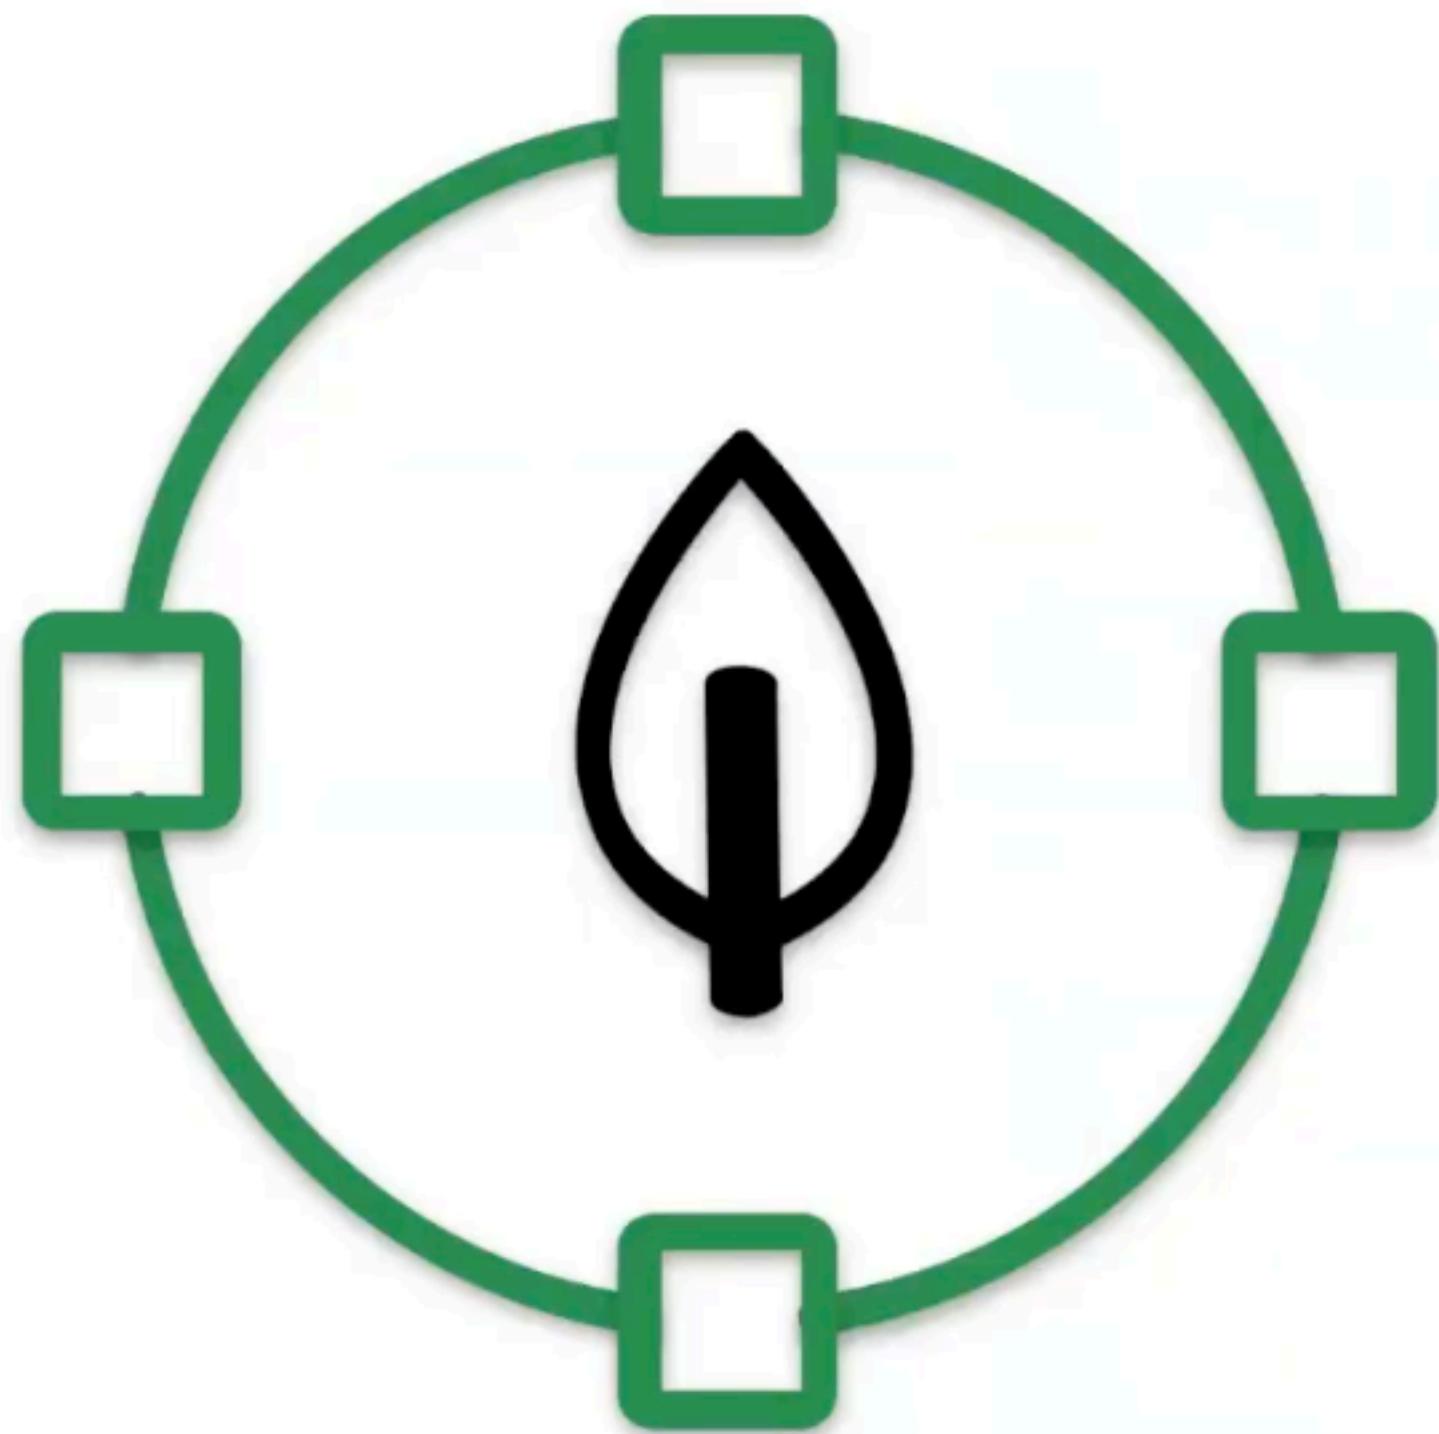
ID	user_id	model	year
20	1	Bentley	1973
21	1	Rolls Royce	1965

MongoDB

```
{  
    first_name: "Paul",  
    surname: "Miller",  
    cell: "447557505611",  
    city: "London",  
    location: [45.123,47.232],  
    profession: ["banking", "finance", "trader"],  
    cars: [  
        {  
            model: "Bentley",  
            year: 1973  
        },  
        {  
            model: "Rolls Royce",  
            year: 1965  
        }  
    ]  
}
```

Embedding vs referencing

Embedding - det vi gjort hittills i våra modeller



```
{  
  _id : ObjectId('AAA'),  
  name: 'Kate Monster',  
  ssn: '123-456-7890',  
  addresses: [  
    {  
      street: '123 Sesame St',  
      city: 'Anytown', cc: 'USA'  
    },  
    {  
      street: '123 Avenue Q',  
      city: 'New York',  
      cc: 'USA'  
    }  
  ]  
}
```

Fristående
objekt,
ingen egen
entitet.

\$lookup

"Left outer join"

- <https://docs.mongodb.com/manual/reference/operator/aggregation/lookup/>

```
{  
  $lookup:  
    {  
      from: <collection to join>,  
      localField: <field from the input documents>,  
      foreignField: <field from the documents of the "from" collection>,  
      as: <output array field>  
    }  
}
```

Embedding

Pro



Retrieve all data with a single query



Avoids expensive JOINs or \$lookup



Update all data with a single atomic operation

Embedding

Con

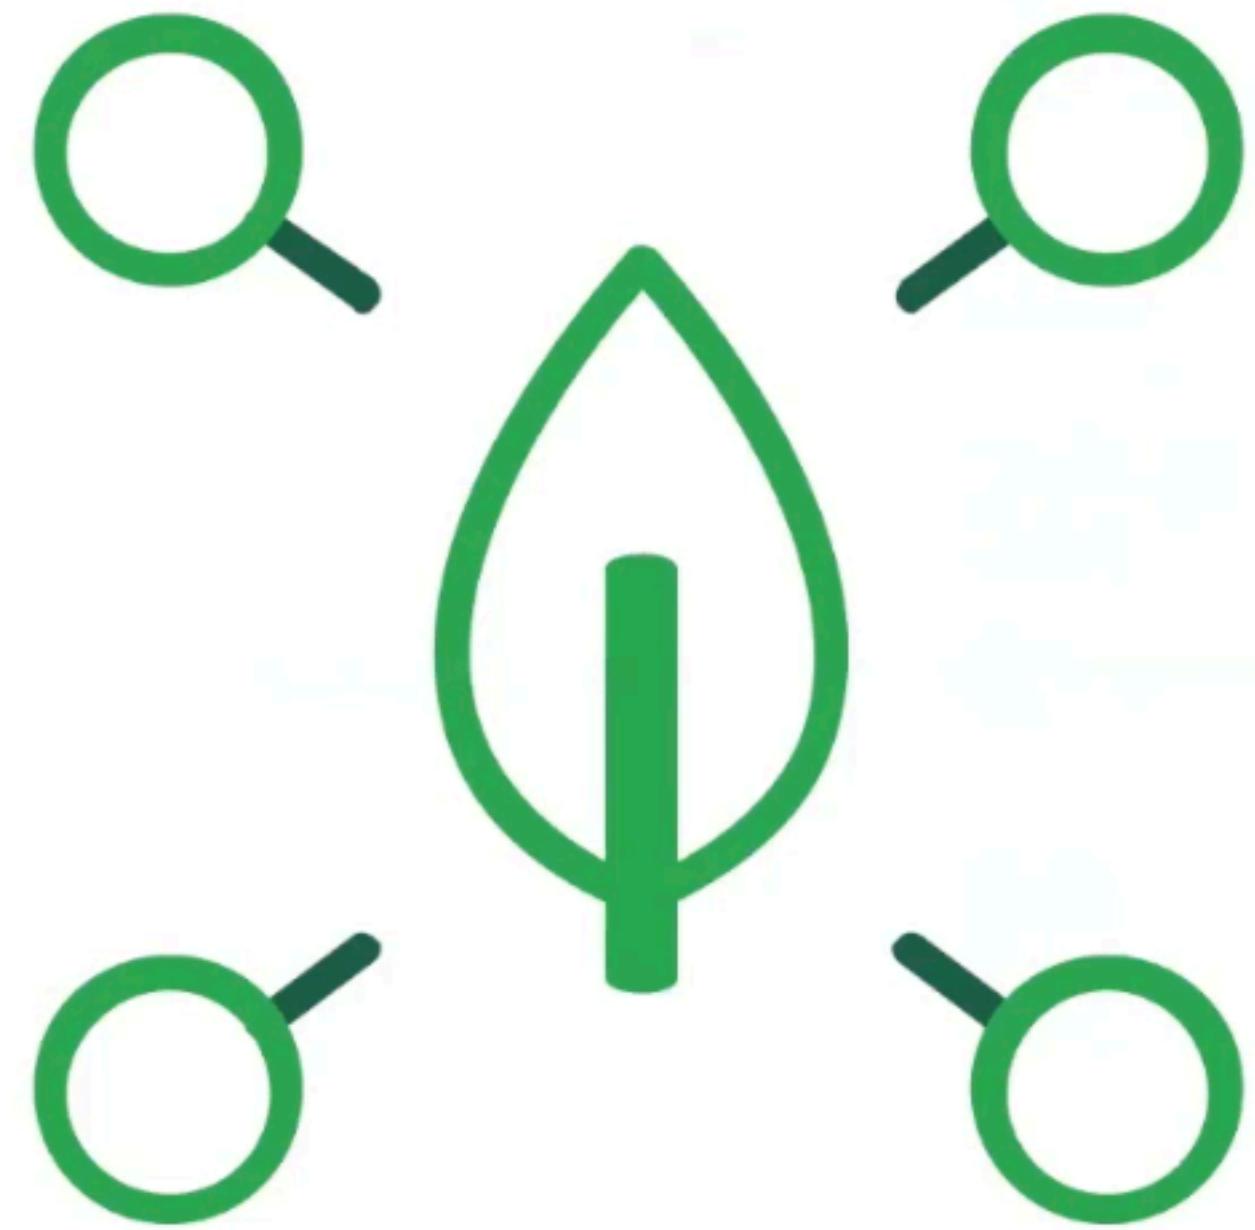


Large docs === more overhead



16 MB Document size limit

Referencing



Referencing

```
{  
    name : 'left-handed smoke shifter',  
    manufacturer : 'Acme Corp',  
    catalog_number: 1234,  
    parts : [  
        ObjectId('AAAA'),  
        ObjectId('BBBB'),  
    ]  
}  
  
{  
    _id : ObjectId('AAAA'),  
    partno : '123-aff-456',  
    name : '#4 grommet',  
    qty: 94,  
    cost: 0.94,  
    price: 3.99  
}  
  
{  
    _id : ObjectId('BBB'),  
    partno : '425-EFF-123',  
    name : '#8 Frombet',  
    qty: 13,  
    cost: 0.34,  
    price: 7.99  
}
```

Inte helt olikt
främmmande
nycklar

Referencing

Pro



Smaller documents



Less likely to reach 16 MB limit



No duplication of data



Infrequently accessed data
not accessed on every query

Referencing

Con



Two queries or \$lookup
required to retrieve all data

Embedding vs referencing

- Här har vi lagt in författaren som ett textvärde, vilket kan vara okej, beroende på vad vi behöver för data i vår applikation.

```
{  
  "_id": {  
    "$oid": "601d1f6104eae3ec898d97b5"  
  },  
  "title": "The Hobbit: An Unexpected Journey",  
  "writer": "J.R.R. Tolkeint",  
  "year": 2012,  
  "franchise": "The Hobbit",  
  "synopsis": "A reluctant hobbit, Bilbo Baggins, sets out to  
}
```

Embedding vs referencing

Embedding

- Vi skulle också kunna tänka oss att embedda ett fristående objekt.

```
{  
  "_id": {  
    "$oid": "601d1f6104eae3ec898d97b5"  
  },  
  "title": "The Hobbit: An Unexpected Journey",  
  "writer": {  
    "name": "J.R.R. Tolkein",  
    "age": 134,  
    "kids": 3  
  },  
  "year": 2012,  
  "franchise": "The Hobbit",  
  "synopsis": "A reluctant hobbit, Bilbo Baggins, sets out to the Lonely Mountain with a spirited group of dwarves to reclaim their mountain home - and the gold within it - from the dragon Smaug."  
}
```

Embedding vs referencing

Referencing

- Eller så skapar vi en Writer-collection, skapar ett dokument där och refererar in det i vårt Movie-dokument.

```
{  
  "_id": {  
    "$oid": "601d1f6104eae3ec898d97b5"  
  },  
  "title": "The Hobbit: An Unexpected Journey",  
  "writer": ObjectId('12345457'),  
  "year": 2012,  
  "franchise": "The Hobbit",  
  "synopsis": "A reluctant hobbit, Bilbo Baggins,  
  dwarves to reclaim their mountain home - and the  
  ring."}
```

```
{  
  "_id": {  
    "$oid": "12345457"  
  },  
  "name": "J.R.R. Tolkein",  
  "age": 134,  
  "kids": 3}
```

Referencing

- Hur vet vi vilket id vi vill spara?
 - Både i relationsdatabaser och i dokumentdatabaser behöver vi hålla reda på vilket id som ska in som ”främmande nyckel”.
 - I våra experiment blir det ett manuellt arbete, i våra webbapplikationer kommer vi att kunna hantera detta mer automatiskt på backend-sidan.
 - *Pseudo-kod* (kod som är lite lik riktig kod, men som inte fungerar utan är mer gjord för att visualisera vad som ska hända för våra mänskliga ögon):
 - `let writerId = createWriter("Tolkien").getID(); // Påhittade metoder, men vi förstår tanken`
 - `addMovie("LOTR", writerId);`

Kardinalitet

One to One



Use Key-Value pairs

User:

```
{  
  _id: ObjectId("AAA"),  
  name: "Joe Karlsson",  
  company: "MongoDB",  
  twitter: "@JoeKarlsson1",  
  twitch: "joe_karlsson",  
  tiktok: "joekarlsson",  
  website: "joekarlsson.com"  
}
```

Kardinalitet

One to Few

```
{  
  _id: ObjectId("AAA"),  
  name: "Joe Karlsson",  
  company: "MongoDB",  
  twitter: "@JoeKarlsson1",  
  twitch: "joe_karlsson",  
  tiktok: "joekarlsson",  
  website: "joekarlsson.com",  
  addresses: [  
    { street: "123 Sesame St", city: "Anytown", cc: "USA" },  
    { street: "123 Avenue Q", city: "New York", cc: "USA" }  
  ]  
}
```

Embedding vs referencing

- Joe pratar om några regler som han har kommit fram till och även om det inte är ”branschregler” på samma sätt som t ex normalisering, så känns de vettiga.

Rule 1:

Favor embedding unless
there is a compelling
reason not to.

Embedding vs referencing

- Tex om vi behöver komma åt alla Writers "för sig", inte som en del av Movie-dokumentet.

Rule 2:

Needing to access an object on its own is a compelling reason not to embed it.

Kardinalitet

One to Many

Products:

```
{  
  _id: ObjectId("123"),  
  name: "left-handed smoke shifter",  
  manufacturer: "Acme Corp",  
  catalog_number: 1234,  
  parts: [  
    ObjectId("AAA"),  
    ObjectId("BBB"),  
    ObjectId("CCC"),  
  ]  
}
```

Parts:

```
{  
  _id: ObjectId("AAA"),  
  partno: "123-ABC-456",  
  name: "#4 grommet",  
  qty: 94,  

```

\$lookup

Exempel

- <https://docs.mongodb.com/manual/reference/operator/aggregation/lookup/#examples>

```
db.orders.insert([
  { "_id" : 1, "item" : "almonds", "price" : 12, "quantity" : 2 },
  { "_id" : 2, "item" : "pecans", "price" : 20, "quantity" : 1 },
  { "_id" : 3 }
])

db.inventory.insert([
  { "_id" : 1, "sku" : "almonds", "description": "product 1", "instock" : 120 },
  { "_id" : 2, "sku" : "bread", "description": "product 2", "instock" : 80 },
  { "_id" : 3, "sku" : "cashews", "description": "product 3", "instock" : 60 },
  { "_id" : 4, "sku" : "pecans", "description": "product 4", "instock" : 70 },
  { "_id" : 5, "sku": null, "description": "Incomplete" },
  { "_id" : 6 }
])
```

\$lookup

Exempel

- <https://docs.mongodb.com/manual/reference/operator/aggregation/lookup/#examples>

```
db.orders.aggregate([
  {
    $lookup:
    {
      from: "inventory",
      localField: "item",
      foreignField: "sku",
      as: "inventory_docs"
    }
  }
])
```

```
{
  "_id" : 1,
  "item" : "almonds",
  "price" : 12,
  "quantity" : 2,
  "inventory_docs" : [
    { "_id" : 1, "sku" : "almonds", "description" : "product 1", "instock" : 120 }
  ]
}
{
  "_id" : 2,
  "item" : "pecans",
  "price" : 20,
  "quantity" : 1,
  "inventory_docs" : [
    { "_id" : 4, "sku" : "pecans", "description" : "product 4", "instock" : 70 }
  ]
}
```

Embedding vs referencing

Rule 3:

Avoid JOINs and \$lookups if they can be, but don't be afraid if they can provide a better schema design.

Kardinalitet

One to Squillions

Hosts:

```
{  
  _id: ObjectId("AAA"),  
  name: "goofy.example.com",  
  ipaddr: "127.66.66.66",  
}
```

Log Message:

```
{  
  _id: ObjectId("123"),  
  time: ISODate("2014-03-28T09:42:41.382Z"),  
  message: "The CPU is on fire!!!",  
  host: ObjectId("AAA"),  
},  
{  
  _id: ObjectId("456"),  
  time: ISODate("2014-03-28T09:42:41.382Z"),  
  message: "Drive is hosed",  
  host: ObjectId("AAA"),  
}
```

Lite som skillnaden
mellan 1:n mot n:1

Referencing

- Tänk att ha egenskapen tv-shows som innehåller referenser till alla program som någonsin visats på kanalen.
- Vänd på det istället och sätt referensen på tv-programmet.

Rule 4:

Arrays should not
grow without bound.

Kardinalitet

Many to Many

Person:

```
{  
  _id: ObjectId("AAF1"),  
  name: "Joe Karlsson",  
  tasks: [  
    ObjectId("ADF9"),  
    ObjectId("AE02"),  
    ObjectId("ZDF2"),  
  ]  
}
```

Tasks:

```
{  
  _id: ObjectId("ADF9"),  
  description: "Learn MongoDB",  
  due_date: ISODate("2014-03-28T09:42:41.382Z"),  
  owners: [  
    ObjectId("AAF1"),  
  ]  
,  
{  
  _id: ObjectId("AE02"),  
  description: "Write lesson plan",  
  due_date: ISODate("2014-03-28T09:42:41.382Z"),  
  owners: [  
    ObjectId("AAF1"),  
  ]  
,
```

Embedding vs referencing

- Skillnad mot RDBMS.
 - RDBMS: Utgår från hur data ser ut.
 - Dokument-db: Utgår ifrån hur du vill använda data.

Rule 5:

How you model your data depends – entirely – on your particular application's data access patterns.